

8.04 Time Series II

Time Series Forecasting with FB Prophet

- Facebook developed an open sourcing Prophet, a forecasting tool available in Python
- It provides intuitive parameters which are easy to tune
- Even someone who lacks deep expertise in time-series forecasting models can use this to generate meaningful predictions for a variety of business use cases

Time Series Forecasting with FB Prophet

- Producing high quality forecasts is not an easy problem for either machines or for most analysts.
- FB observed two main themes in the practice of creating a variety of business forecasts:
 - Completely automatic forecasting techniques can be fragile and they are often too inflexible to incorporate useful assumptions.
 - Analysts who can produce high quality forecasts are quite rare because forecasting is a specialized data science skill requiring substantial experience.

Key Highlights of FB Prophet

- Very fast, since it's built in Stan, a programming language for statistical inference written in C++
- An additive regression model where non-linear trends are fit with yearly, weekly, and daily seasonality, plus holiday effects
- Robust to missing data and shifts in the trend, and typically handles outliers
- Easy procedure to tweak and adjust forecast while adding domain knowledge or business insights

FB Prophet Forecasting Model

- FB Prophet uses a decomposable time series model with three main model components: trend, seasonality, and holidays.
- They are combined in the following equation:
$$y(t) = g(t) + s(t) + h(t) + \epsilon_t$$
- $g(t)$: piecewise linear or logistic growth curve for modeling non-periodic changes in time series
- $s(t)$: periodic changes (e.g. weekly/yearly seasonality)
- $h(t)$: effects of holidays (user provided) with irregular schedules
- ϵ_t : error term accounts for any unusual changes not accommodated by the model

FB Prophet Forecasting Model

- FB Prophet uses a decomposable time series model with three main model components: trend, seasonality, and holidays.
- They are combined in the following equation:
$$y(t) = g(t) + s(t) + h(t) + \epsilon_t$$
- $g(t)$: piecewise linear or logistic growth curve for modeling non-periodic changes in time series
- $s(t)$: periodic changes (e.g. weekly/yearly seasonality)
- $h(t)$: effects of holidays (user provided) with irregular schedules
- ϵ_t : error term accounts for any unusual changes not accommodated by the model